

Estimating Surface Orientations Based on Monocular Image Queues

Alexander Norton
Computer Science Department,
Colorado State University
norton@cs.colostate.edu

Abstract

Abstract - Recovering the 3 dimensional structure of a scene from a single 2 dimensional video stream is a task that humans have very little trouble with. However, when computers attempt the same process, the products are either very simple, or very prone to error. We present a algorithm that does a simple image segmentation based with the goal of finding the vertical surface orientations. Based upon the relative orientations of the different regions of the image, the relative depth of objects of interest within the video can then be estimated. We provide a quantitative analysis of the algorithm of a set of monocular outdoor images and a qualitative analysis on video data.

1. Introduction

Reconstruction of the 3 dimensional scene geometry is an important step in understanding a scene and interpreting the interactions between the objects within that scene. Understanding the relative depth of two distinct object can help inform the relationship between them. For example, one could infer if the two objects would be able to touch based upon not only how close they are in the pixel space, but also how close they fall in relative depth. There has been work on create a full reconstruction in the form of a pop-up model from a single monocular image [3, 4, 5]. Gupta, Efros and Hebert [2] improved upon this by creating volumetric versions by trying to recreate the blocks in the monocular image using a 3 dimensional model. Other forms of scene reconstruction depend upon used know motion of the camera, stereo vision or other similar techniques to find the geometry.

We propose a algorithm that combines these two efforts. It attempts to reconstruct the scene geometry from a monocular video instead of a monocular image. This allows the algorithm to incorporate temporal informa-



Figure 1: A monocular image with the over-segmentation produced by Felzenszwalb et al. [1]. These segmenst are referred to as super pixels.

tion into the surface estimations. However, structure from motion or similar techniques cannot be applied to the video since there is no information about the camera or defined camera locations. As a result, the algorithm relies upon the methods developed by Hoiem et al. [3] and Sexana et al. [5] to find the vertical surface orientations. From this the relative depth of different regions of the image can be estimated to provide object information. The goal is to create a depth segmentation of the image.

The algorithm uses supervised learning to create a classifier for image regions. The algorithm will start with an over-segmentation of the image to produce super pixels. Super pixels are simply image regions that statistical information can gathered for. The statistical information is used as input for the learned classifier which is used to decide the surface orientations. From the surface orientations we can gather depth information by simply using where the vertical surfaces intersect the horizontal surface defined by the ground plane.

2. Super Pixels

The first step in finding a usable depth segmentation is to perform an over-segmentation of the image. An example of the over-segmentation can be seen in Figure 1. These super pixels will be used as the units of surface. We will construct the surfaces out of the individual su-




Figure 2: One color channel of a source image with the Law's mask convolutions. The average response from these masks is used as one of the features for the classifier.

per pixels and the borders between surfaces will be assumed to fall on the boundaries between super pixel. As a result of this, the only requirement of the over-segmentation is that there is no single super pixel that spans two real world surfaces within the image. If a super pixel does fall on two different real world surfaces, then it would have multiple orientation and automatically cause the depth segmentation to be invalid. The algorithm uses the segmentation algorithm described by Felzenszwalb et al. [1] to produce the super pixel image. Any segmentation algorithm could be used in instead as long as it doesn't produce any segments that contain two distinct real world surfaces.

For each super pixel in the image a set of statistics are calculated. The major distinguishing factors between horizontal and vertical surfaces are the textures and line orientations on the surface. As a result, the texture is calculated as the response to a set of 3x3 Law's masks for the different color channels of the image. The Law's masks are created using the three vectors $[1, 2, 1]$, $[1, 0, -1]$, and $[1, -1, 1]$. These vectors are multiplied to get 9 3x3 masks that are convolved with the color channel. This produces a total of 27 responses from the Law's masks. A source image and some example responses for the Law's masks can be seen in Figure 2. A histogram of gradient orientations is calculated for each super pixel as well. 18 buckets are used for the histogram of gradients, and each color channel uses a different histogram. This gives a total of 54 buckets from gradient histograms. This gives an 81 dimensional space that is used to represent the super pixel in the learned classifier.

3. Learned Classifier

The collection of features used for each super pixel is fed into a learning classifier that will inform the algorithm about the orientations of the super pixels. Both grouping the super pixels together into larger image regions and testing super pixels completely independently was tried. Two different learned classifiers were used. The discrete form of Adaboost and a simple k-nearest

classifier were both tested and the results compared. These were used because of the simplicity involved in using them correctly.

Both individual super pixels and small image regions were used as input for the classifier. These were tested individually for the training and testing phases. The rationale behind testing image region is that the orientation can be informed by what lies around the super pixel. For example, take the case of a brick wall where the boundaries between the super pixels would fall at the edge of each brick. If we include adjacent super pixels in the calculation of the score for the relevant super pixel, then we can include more of the texture of the brick wall. We are pulling more information into the surface estimations in the hopes that we will get a better response.

The classifier was tested on 17 labeled training images. These images were labeled by hand with the possibilities of vertical, horizontal, sky or other. For the classifier, the category of sky was changed for vertical. There is reason to find sky separately from other image categories since it can be interpreted as having an infinite depth. However, to simplify the classifier, this category was treated as vertical.

Once an image has been over-segmented, each super pixel is used as input into the classifier. The results of the classifier are then used to directly label the image as either vertical or horizontal. The image is split into these two separate segments as in Figure 3.

4. Results

The algorithm was tested on a set of 300 images of outdoor scenes. Of these images, the number of super pixel correctly labels was used to determine if the algorithm adequately labeled that image. The k-nearest classifier was not rigorously tested as it was simply a proof of concept. However, based upon some qualitative analysis, the k-nearest classifier performed much better than expected. It would correctly label enough super pixels to identify the vertical and horizontal images.




Figure 3: The output of the classifier. Image regions of the source image (the top image) are input into the classifier. The super pixels that the classifier decided were horizontal are on the left. The vertical are on the right.

For the Adaboost classifier, both individual super pixels and image regions were tested. When building the image regions for training, the region for a particular super pixel included all adjacent super pixels that were of the same class. This would mean that the training didn't include a mixing of classes and would avoid confusing the classifier. For testing, the region simply consisted of all adjacent super pixels. As we do not know what class two super pixel belong to during testing, we assume that the most common case (i.e. they do belong to the same surface) is the correct case.

For individual super pixels with the Adaboost classifier, the algorithm correctly estimated the horizontal and vertical surfaces for 145 of the 300 test images. For regions of super pixels with the Adaboost classifier, the algorithm got 104 image correctly estimated. This is a accuracy of %48.3 for individual super pixels and %34.6 for the super pixels that were grouped into regions. It is important to note that while there was significant overlap between the sets of images that both correctly estimated, there were several notable exceptions that did not appear in both sets of correctly guessed images.

Qualatatively, on the limited number of videos of

outdoor scenes that the algorithm has been run on, it has done well segmenting the image. Interestingly the results of the algorithm from one frame to the next can change drastically. This would indicate that small changes to the underlying pixels can result in a very different response from classifier. This may be due to a drastic change to the segmentation.

5. Conclusion

We set the goal of creating an algorithm that could perform an image segmantion that would allow for estimations of horizontal and vertical surface orientation estimation. We have taken important steps towards accomplishing this goal. Given a relatively simple outdoor scene, we can correctly determine what surfaces within the image are horizontal and which are vertical. From the succesfull image segmenations we have shown that simple texture and gradient orientations contain enough information to estimate the 3 dimensional orientation of the underlying surface.

However, we have also shown that something that should logically improve the surface orientation estimation, namely taking the neighbors of a super pixel into



Figure 4: Examples of correct segmented images. The first column is the original source image. The middle column is the horizontal segment of the image. The right column is the vertical segment of the image

account, actually decreased the quality of the vertical surface orientation estimation. Hoiem et al. [3] use a learning pairing function to guess when two super pixels belong to the same class. Also the set of images correctly labeled by region classifier is not a subset of the images correctly labeled by the individual classifier. These would indicate that the performance of the algorithm could possibly be improved by not necessarily including every neighbor when creating the region. Something similar to the learned pairing function used in Hoiem et al. [3] should be explored.

References

- [1] P. F. Felzenszwalb and D. P. Huttenlocher. Efficient graph-based image segmentation. *Int. J. Comput. Vision*, 59(2):167–181, Sept. 2004.
- [2] A. Gupta, A. A. Efros, and M. Hebert. Blocks world revisited: Image understanding using qualitative geometry and mechanics. In *European Conference on Computer Vision (ECCV)*, 2010.
- [3] D. Hoiem, A. A. Efros, and M. Hebert. Geometric context from a single image. In *Proceedings of the Tenth IEEE International Conference on Computer Vision (ICCV'05) Volume 1 - Volume 01*, ICCV '05, pages 654–661, Washington, DC, USA, 2005. IEEE Computer Society.
- [4] D. Hoiem, A. A. Efros, and M. Hebert. Closing the loop on scene interpretation. In *Proc. Computer Vision and Pattern Recognition (CVPR)*, June 2008.
- [5] A. Saxena, M. Sun, and A. Y. Ng. Make3d: Learning 3d scene structure from a single still image. *IEEE Trans. Pattern Anal. Mach. Intell.*, 31(5):824–840, May 2009.